

▼ SUPPORT VECTOR MACHINE ALGORITHM on Amazon Fine Food Reviews

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (Rating of 4 or 5) or negative (rating of 1 or 2).

▼ Loading,Cleaning & Preprocessing the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all

3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

▼ Import all Required Libraries

```
1 %matplotlib inline
2 import warnings
3 import sqlite3
4 import pandas as pd
5 import numpy as np
6 import nltk
7 import string
8 import matplotlib.pyplot as plt
9 import seaborn as sns
10 from sklearn.feature_extraction.text import TfidfTransformer
11 from sklearn.feature_extraction.text import TfidfVectorizer
12 from sklearn.feature_extraction.text import CountVectorizer
13 from sklearn.metrics import confusion_matrix
14 from sklearn import metrics
```

```

15 from sklearn.metrics import roc_curve, auc
16 from nltk.stem.porter import PorterStemmer
17
18 import re
19 import string
20 from nltk.corpus import stopwords
21 from nltk.stem import PorterStemmer
22 from nltk.stem.wordnet import WordNetLemmatizer
23
24 from gensim.models import Word2Vec
25 from gensim.models import KeyedVectors
26 import pickle
27
28 from tqdm import tqdm
29 import os
30
31 warnings.filterwarnings("ignore")

```

↳ /usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated
import pandas.util.testing as tm

▼ Pull the dataset from Google Drive & mount

```

1 from google.colab import drive
2 drive.mount('/content/drive')

```

↳ Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pf

Enter your authorization code:

.....

Mounted at /content/drive

▼ Connect to Database and execute the query

```

1 con = sqlite3.connect("/content/drive/My Drive/Colab Notebooks/database.sqlite")
2
3 filtered_data=pd.read_sql_query("""SELECT * FROM Reviews WHERE Score != 3""",con);
4 filtered_data.head(3)

```

↳

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	

▼ Data Cleaning & Preparation

```

1 filtered_data.shape
2
3 def partition(x):
4     if x < 3 :
5         return 'negative'
6     return 'positive'
7

```

☞ Number of datapoints (525814, 10)

```
1 display = pd.read_sql_query("""
2 SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
3 FROM Reviews
4 GROUP BY UserId
5 HAVING COUNT(*)>1
6 """, con)
```

☞ (80668, 7)

```
1 display= pd.read_sql_query("""
2 SELECT *
3 FROM Reviews
4 WHERE Score != 3 AND UserId="AR5J8UI46CURR"
5 ORDER BY ProductID
6 """, con)
7 display.head()
```

3/32

```

1 sorted_data=filtered_data.sort_values('ProductId',axis=0,ascending=True,inplace=False,kind='quicksort',na_positior
2
3 final_data=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"},keep='first',inplace=False)
4 final_data.shape

```

```
↳ (364173, 10)
```

```
1 (final_data['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
↳ 69.25890143662969
```

```

1 display= pd.read_sql_query("""
2 SELECT *
3 FROM Reviews
4 WHERE Score != 3 AND Id=44737 OR Id=64422
5 ORDER BY ProductID
6 """, con)
7
8 display.head()

```

```
↳
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	1	5 1
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	2	4 1

```
1 final_data=final_data[final_data.HelpfulnessNumerator<=final_data.HelpfulnessDenominator]
```

```
1 nltk.download('stopwords')
```

```

↳ [nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
True

```

```

1 stopping_words = set(stopwords.words('english'))
2 print(stopping_words)

```

```
↳ {'itself', 'against', 'here', "didn't", 'myself', 'under', "you're", 'most', 'no', 'all', "should've", 'will', '
```

```

1 def clean_html(text):
2     clean_r = re.compile('<,*?>')
3     clean_text = re.sub(clean_r,'',text)
4     return clean_text
5
6 def Clean_punc(text):
7     clean_sentence = re.sub(r'[?|!|\'|\"|#]',r' ',text)
8     clean_data = re.sub(r'[.,|)|(|\|/)]',r' ',clean_sentence)
9     return clean_data

```

```

1 from tqdm import tqdm
2 import os
3 import pdb
4 import pickle
5
6 from tqdm import tqdm
7 import os
8 import pdb
9 import pickle
10

```

```

11 stem_no = nltk.stem.SnowballStemmer('english')
12
13 if not os.path.isfile('final_data.sqlite'):
14     final_string=[]
15     all_positive_words=[]
16     all_negative_words=[]
17     for i,sentence in enumerate(tqdm(final_data['Text'].values)):
18         filtered_sentence=[]
19         sent_without_html_tags=clean_html(sentence)
20         #pdb.set_trace()
21         for w in sent_without_html_tags.split():
22             for cleaned_words in Clean_punc(w).split():
23                 if ((cleaned_words.isalpha()) & (len(cleaned_words) > 2)):
24                     if(cleaned_words.lower() not in stopping_words) :
25                         stemming=(stem_no.stem(cleaned_words.lower())).encode('utf8')
26                         filtered_sentence.append(stemming)
27                         if(final_data['Score'].values[i]=='positive':
28                             all_positive_words.append(stemming)
29                         if(final_data['Score'].values[i]=='negative':
30                             all_negative_words.append(stemming)
31         str1 = b" ".join(filtered_sentence)
32         final_string.append(str1)
33
34     final_data['Cleaned_text']=final_string
35     final_data['Cleaned_text']=final_data['Cleaned_text'].str.decode("utf-8")
36
37     conn = sqlite3.connect('final_data.sqlite')
38     cursor=conn.cursor
39     conn.text_factory = str
40     final_data.to_sql('Reviews',conn,schema=None,if_exists='replace',index=True,index_label=None,chunksize=None,dt
41     conn.close()
42
43
44     with open('positive_words.pkl','wb') as f :
45         pickle.dump(all_positive_words,f)
46     with open('negative_words.pkl','wb') as f :
47         pickle.dump(all_negative_words,f)

```

100%|██████████| 364171/364171 [06:07<00:00, 991.13it/s]

```
1 final_data['total_words'] = [len(x.split()) for x in final_data['Cleaned_text'].tolist()]
```

```
1 final_data.sort_values(by=['Time'], inplace=True, ascending=True)
```

```

1 final_data_100K=final_data[0:100000]
2 amazon_polarity_labels=final_data_100K['Score'].values
3 final_data_100K.head(2)

```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score
138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	0	positive
138683	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	2	2	positive

```
1 final_data_40K = final_data[0:30000]
2 amazon_polarity_labels_40K = final_data_40K['Score'].values
```

▼ Split the dataset in Train , Test & Cross Validation

```
1 from sklearn.model_selection import train_test_split
2 from sklearn.neighbors import KNeighborsClassifier
3 from sklearn.metrics import accuracy_score
4 from sklearn.model_selection import cross_val_score
5 from collections import Counter
6 from sklearn.metrics import confusion_matrix
7 from sklearn.metrics import classification_report
8
9 X_1,X_Test,Y_1,Y_Test = train_test_split(final_data_100K,amazon_polarity_labels,test_size=0.2,random_state=0)
10 X_Train,X_CV,Y_Train,Y_CV = train_test_split(X_1,Y_1,test_size=0.2)
```

▼ APPLY BAG OF WORDS VECTORIZATION TECHNIQUE USING LINEAR SVM (SGDCLASSIF FIND THE BEST ALPHA

```
1 print(X_Train.shape, Y_Train.shape)
2 print(X_CV.shape, Y_CV.shape)
3 print(X_Test.shape, Y_Test.shape)
4
5 print("="*100)
6
7
8 count_vector=CountVectorizer(min_df=1)
9 X_Train_data_bow=(count_vector.fit_transform(X_Train['Cleaned_text']).values)
10 X_Test_data_bow=(count_vector.transform(X_Test['Cleaned_text']).values)
11 X_CV_data_bow=(count_vector.transform(X_CV['Cleaned_text']).values)
12
13 print("After vectorizations")
14 print(X_Train_data_bow.shape, Y_Train.shape)
15 print(X_CV_data_bow.shape, Y_CV.shape)
16 print(X_Test_data_bow.shape, Y_Test.shape)
17 print("="*100)
```

```
↳ (64000, 12) (64000,)
   (16000, 12) (16000,)
   (20000, 12) (20000,)
=====
After vectorizations
(64000, 29463) (64000,)
(16000, 29463) (16000,)
(20000, 29463) (20000,)
=====
```

```
1 from sklearn.model_selection import GridSearchCV
2 from scipy.stats import randint as sp_randint
3 from sklearn.model_selection import cross_val_score
4 from sklearn.linear_model import SGDClassifier
5
6 def SVM_Linear_SVM(x_training_data,y_training_data,regularization):
7     grid_params = { 'alpha' : [10**x for x in range(-4,5)]
8                     }
9     LinearSVM = SGDClassifier(loss='hinge',penalty=regularization,random_state=None, class_weight=None)
10    clf=GridSearchCV(LinearSVM,grid_params,scoring='roc_auc',return_train_score=True,cv=10)
11    clf.fit(x_training_data,y_training_data)
12    results = pd.DataFrame.from_dict(clf.cv_results_)
13    results = results.sort_values(['param_alpha'])
```

```

14 train_auc= results['mean_train_score']
15 train_auc_std= results['std_train_score']
16 cv_auc = results['mean_test_score']
17 cv_auc_std= results['std_test_score']
18 best_alpha = results['param_alpha']
19 #print(type(alpha))
20 #print(alpha)
21 #log_c=np.log10(list(results["param_C"]))
22 print(clf.best_score_)
23 print(clf.best_params_)
24 plt.plot(best_alpha, train_auc, label='Train AUC')
25 plt.plot(best_alpha, cv_auc, label='CV AUC')
26 plt.scatter(best_alpha, train_auc, label='Train AUC points')
27 plt.scatter(best_alpha, cv_auc, label='CV AUC points')
28 plt.legend()
29 plt.xlabel("Alpha : hyperparameter")
30 plt.ylabel("AUC")
31 plt.title("Hyper parameter Vs AUC plot")
32 plt.grid()
33 plt.show()
34 return results,clf,LinearSVM

```

▼ Call Linear SVM using L2 Regularization

```


1 results,best_alpha,linear_svm = SVM_Linear_SVM(X_Train_data_bow,Y_Train,'l2')
2 results.head()

```

```

0.9266877065638912
{'alpha': 0.001}

```



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_score	split1
0	1.228909	0.056521	0.016896	0.001494	1e-05	{'alpha': 1e-05}	0.905648	
1	0.641934	0.049558	0.016620	0.000650	0.0001	{'alpha': 0.0001}	0.923277	
2	0.377793	0.019492	0.016994	0.001798	0.001	{'alpha': 0.001}	0.930186	
3	0.297809	0.009442	0.016568	0.000745	0.01	{'alpha': 0.01}	0.907816	
4	0.297083	0.006493	0.017016	0.001693	0.1	{'alpha': 0.1}	0.704151	

▼ Got Best Alpha = 0.001 and AUC SCORE = 0.9266 ON TRAINING DATA

```

1 def Find_best_Alpha(best_alpha) :
2     best_Alpha = best_alpha.best_params_
3     best_alpha=best_Alpha.get("alpha")
4     print(best_alpha)
5     return best_alpha

```

```
1 best_alpha = Find_best_Alpha(best_alpha)
```


↗ 0.001

```

1 from sklearn.metrics import roc_curve, auc
2 from sklearn.calibration import CalibratedClassifierCV
3
4 linear_svm = SGDClassifier(loss='hinge',penalty='l2',random_state=None, class_weight=None,alpha=best_alpha)
5 clf=linear_svm.fit(X_Train_data_bow,Y_Train)
6 calibrated_model=CalibratedClassifierCV(clf,cv='prefit')
7 model_m=calibrated_model.fit(X_Train_data_bow,Y_Train)
8 pred_test_data=linear_svm.predict(X_Test_data_bow)
9 y_train_predicted_prob = model_m.predict_proba(X_Train_data_bow)[: ,1]
10 y_test_predicted_prob=model_m.predict_proba(X_Test_data_bow)[: ,1]
11 train_fpr, train_tpr, train_thresholds=roc_curve(Y_Train,y_train_predicted_prob,pos_label='positive')
12 test_fpr, test_tpr, test_thresholds = roc_curve(Y_Test, y_test_predicted_prob,pos_label='positive')
13 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
14 plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
15 plt.legend()
16 plt.xlabel("Alpha: hyperparameter")
17 plt.ylabel("AUC")
18 plt.title("ERROR PLOTS")
19 plt.grid()
20 plt.show()

```

↗



▼ ON TEST DATA WE GOT AUC AS 0.9280 FOR ALPHA = 0.001

```

1 feature_names = count_vector.get_feature_names()
2 print(feature_names)
3 coefs_with_fns = sorted(zip(linear_svm.coef_[0], feature_names))
4 top_features = zip(coefs_with_fns[:20], coefs_with_fns[:-(20 + 1):-1])
5 list(top_features)

```

↗


```
[ 'aaa', 'aaaaaaaaagghh', 'aaaaah', 'aaaaahhhhhhhhhhhhhhh', 'aaah', 'aaahhhhh', 'aachen', 'aacur', 'aadp', 'aaf',
[(-1.058942800585861, 'worst'), (0.5322610283239771, 'delici'),
((-0.7799799974810454, 'disappoint'), (0.4731209140657581, 'excel')),
((-0.7665897829320187, 'terribl'), (0.4641941043664035, 'amaz')),
((-0.7152606271607351, 'horribl'), (0.45303559224220846, 'best')),
((-0.7074496686738002, 'threw'), (0.4374136752683411, 'perfect')),
((-0.698522858974446, 'aw'), (0.4240234607193107, 'great')),
((-0.6873643468502516, 'return'), (0.3648833464610924, 'addict')),
((-0.6159498692554216, 'tasteless'), (0.3570723879741582, 'nice')),
((-0.6025596547063876, 'bland'), (0.3548406855493194, 'awesom')),
((-0.5858218865201033, 'stale'), (0.3314078100885153, 'satisfi')),
((-0.5032488968010768, 'unfortun'), (0.32582855402641864, 'wonder')),
((-0.49990134316381896, 'disgust'), (0.32359685160158047, 'yummi')),
((-0.4932062358893016, 'poor'), (0.31578589311464605, 'smooth')),
((-0.482047723765109, 'money'), (0.3113224882649689, 'beat')),
((-0.4697733604285002, 'stuck'), (0.3057432322028735, 'happi')),
((-0.4574989970918883, 'yuck'), (0.30239567856561533, 'hook')),
((-0.4563831458794691, 'gross'), (0.30239567856561467, 'tasti')),
((-0.4552672946670491, 'nasti'), (0.29904812492835753, 'rich')),
((-0.4362978240559228, 'ined'), (0.2990481249283572, 'uniqu')),
((-0.4362978240559222, 'weak'), (0.2945847200786804, 'fantast'))]
```

▼ TOP 20 POSITIVE & NEGATIVE FEATURES

```
1 from sklearn.metrics import roc_auc_score
2
3 roc_auc_score(Y_Test,y_test_predicted_prob)
```

```
0.9280531485969552
```

```
1 from sklearn.metrics import classification_report,confusion_matrix
2
3 print(classification_report(Y_Test,pred_test_data))
4 print(confusion_matrix(Y_Test,pred_test_data))
```

```
precision    recall  f1-score   support

negative     0.82     0.44     0.57      2457
positive     0.93     0.99     0.96     17543

accuracy                   0.92      20000
macro avg     0.87     0.71     0.76      20000
weighted avg   0.91     0.92     0.91      20000

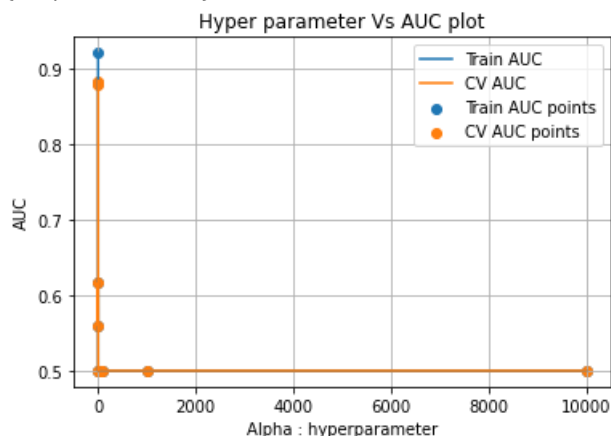
[[ 1069  1388]
 [  237 17306]]
```

▼ SVM on Amazon Fine Food Review using L1 Regularization

```
1 results,best_alpha,linear_svm = SVM_Linear_SVM(X_Train_data_bow,Y_Train,'l1')
2 results.head()
```

0.8839167307373945

{ 'alpha': 0.0001}



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_score	split1
0	4.734316	0.663837	0.019478	0.000709	0.0001	{'alpha': 0.0001}	0.893648	
1	0.740307	0.322096	0.019451	0.000378	0.001	{'alpha': 0.001}	0.881980	
2	0.369847	0.158847	0.020437	0.002277	0.01	{'alpha': 0.01}	0.592772	
3	0.740072	0.042864	0.019065	0.000328	0.1	{'alpha': 0.1}	0.596985	
4	0.413893	0.038416	0.018740	0.000523	1	{'alpha': 1}	0.500000	

▼ TF-IDF Vectorization Technique on SUPPORT VECTOR MACHINE

```

1 from sklearn.feature_extraction.text import TfidfVectorizer
2
3 print(X_Train.shape, Y_Train.shape)
4 print(X_CV.shape, Y_CV.shape)
5 print(X_Test.shape, Y_Test.shape)
6
7 print("="*100)
8
9
10 tfidf_vector=TfidfVectorizer(min_df=10)
11 X_Train_data_tfidf=(tfidf_vector.fit_transform(X_Train['Cleaned_text'].values))
12 X_Test_data_tfidf=(tfidf_vector.transform(X_Test['Cleaned_text'].values))
13 X_CV_data_tfidf=(tfidf_vector.transform(X_CV['Cleaned_text'].values))
14
15 print("After vectorizations")
16 print(X_Train_data_tfidf.shape, Y_Train.shape)
17 print(X_CV_data_tfidf.shape, Y_CV.shape)
18 print(X_Test_data_tfidf.shape, Y_Test.shape)
19 print("="*100)

```




```
(64000, 12) (64000,)
(16000, 12) (16000,)
(20000, 12) (20000,)
```

```
=====
After vectorizations
```

```
(64000, 6775) (64000,)
(16000, 6775) (16000,)
(20000, 6775) (20000,)
=====
```

```
1 results,best_alpha,linear_svm = SVM_Linear_SVM(X_Train_data_tfidf,Y_Train,'l2')
2 results.head()
```

```
0.9353531913390751
{'alpha': 0.0001}
```



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_score	split1
0	0.200958	0.006349	0.016475	0.001642	0.0001	{'alpha': 0.0001}	0.939161	
1	0.167422	0.002496	0.017399	0.002993	0.001	{'alpha': 0.001}	0.934442	
2	0.151980	0.004264	0.015957	0.000469	0.01	{'alpha': 0.01}	0.934696	
3	0.360621	0.023838	0.015771	0.000320	0.1	{'alpha': 0.1}	0.654612	
4	0.256873	0.003282	0.016171	0.001593	1	{'alpha': 1}	0.629172	

▼ BEST ALPHA = 0.0001 & AUC SCORE = 0.9353 ON TRAINING DATA

```
1 best_alpha = Find_best_Alpha(best_alpha)
```


```
0.0001
```

```
1 linear_svm = SGDClassifier(loss='hinge',penalty='l2',random_state=None, class_weight=None,alpha=best_alpha)
2 clf=linear_svm.fit(X_Train_data_tfidf,Y_Train)
3 calibrated_model=CalibratedClassifierCV(clf,cv='prefit')
4 model_m=calibrated_model.fit(X_Train_data_tfidf,Y_Train)
5 pred_test_data=linear_svm.predict(X_Test_data_tfidf)
6 y_train_predicted_prob = model_m.predict_proba(X_Train_data_tfidf)[:,-1]
7 y_test_predicted_prob=model_m.predict_proba(X_Test_data_tfidf)[:,-1]
8 train_fpr, train_tpr, train_thresholds=roc_curve(Y_Train,y_train_predicted_prob,pos_label='positive')
9 test_fpr, test_tpr, test_thresholds = roc_curve(Y_Test, y_test_predicted_prob,pos_label='positive')
10 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
```

```

11 plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
12 plt.legend()
13 plt.xlabel("Alpha: hyperparameter")
14 plt.ylabel("AUC")
15 plt.title("ERROR PLOTS")
16 plt.grid()
17 plt.show()

```



```

1 feature_names = tfidf_vector.get_feature_names()
2 print(feature_names)
3 coefs_with_fns = sorted(zip(linear_svm.coef_[0], feature_names))
4 top_features = zip(coefs_with_fns[:20], coefs_with_fns[:-(20 + 1):-1])
5 list(top_features)

```

```

['abandon', 'abc', 'abdomin', 'abil', 'abl', 'abroad', 'absenc', 'absent', 'absolut', 'absolutley', 'absorb', 'a
[(-3.5252305591227597, 'worst'), (3.1691492148895763, 'great')),
((-3.4005030003873884, 'disappoint'), (2.6431849273213213, 'best')),
((-3.011837811348974, 'return'), (2.346785762435066, 'love')),
((-2.8599641269962075, 'horribl'), (2.2155651378079484, 'delici')),
((-2.8490643379944345, 'terribl'), (1.973408817473574, 'good')),
((-2.6997285069022863, 'aw'), (1.8981549152660266, 'excel')),
((-2.551410076986622, 'threw'), (1.8505124817553666, 'perfect')),
((-2.2322630718087475, 'wast'), (1.8235615796747322, 'nice')),
((-2.222731395767679, 'stale'), (1.5899723194087558, 'wonder')),
((-2.1713121701828735, 'money'), (1.5875181466524377, 'amaz')),
((-2.1643257959150217, 'bland'), (1.554797524815107, 'favorit')),
((-2.0874293899363803, 'tasteless'), (1.3285533753518113, 'find')),
((-1.8424766374606494, 'refund'), (1.2952990206015922, 'awesom')),
((-1.8317515775148232, 'disgust'), (1.2663691707134266, 'tasti')),
((-1.8199214620059139, 'poor'), (1.2245791948877016, 'addict')),
((-1.7322759294298344, 'stuck'), (1.193569113054718, 'smooth')),
((-1.7126635922215447, 'weak'), (1.1820805784546944, 'satisfi')),
((-1.6849467359259969, 'unfortun'), (1.148515450544414, 'fast')),
((-1.5951771543881406, 'gross'), (1.1385650966912313, 'easi')),
((-1.564336208325889, 'bad'), (1.1284937586979071, 'yummi'))]

```

```
1 roc_auc_score(Y_Test,y_test_predicted_prob)
```

```
0.9362961422472338
```

```

1 print(classification_report(Y_Test,pred_test_data))
2 print(confusion_matrix(Y_Test,pred_test_data))

```

	precision	recall	f1-score	support
negative	0.90	0.33	0.48	2457
positive	0.91	0.99	0.95	17543
accuracy			0.91	20000
macro avg	0.91	0.66	0.72	20000
weighted avg	0.91	0.91	0.89	20000


```
[[ 804 1653]
 [ 90 17453]]
```

AND FOR BEST ALPHA = 0.0001 , AUC SCORE = 0.9326

▼ SVM on Amazon Fine Food Review using L1 Regularization

```
1 results,best_alpha,linear_svm = SVM_Linear_SVM(X_Train_data_tfidf,Y_Train,'l1')
2 results.head()
```

```
0.9180460072517856
{'alpha': 0.0001}
```



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_score	split1
0	0.282197	0.004999	0.016205	0.000478	0.0001	{'alpha': 0.0001}	0.921746	
1	0.201957	0.009141	0.016437	0.001334	0.001	{'alpha': 0.001}	0.690765	
2	0.192060	0.007254	0.015663	0.000876	0.01	{'alpha': 0.01}	0.500000	
3	1.452912	0.292737	0.014967	0.000299	0.1	{'alpha': 0.1}	0.500000	
4	0.337768	0.006949	0.016103	0.003331	1	{'alpha': 1}	0.500000	

▼ Avg Word2Vec Vectorization Technique on SUPPORT VECTOR MACHINE Algorithm

```
1 from gensim.models import Word2Vec
2 from gensim.models import KeyedVectors
3 import pickle
4
5 list_of_sent_train_avgw2v=[]
```

```

6 list_of_sent_test_avg2v=[]
7 list_of_sent_cv_avg2v=[]
8 for sent_train_avg2v in tqdm(X_Train['Cleaned_text'].values):
9     list_of_sent_train_avg2v.append(sent_train_avg2v.split())

```

100% |██████████| 64000/64000 [00:00<00:00, 78995.77it/s]

```

1 for sent_test_avg2v in tqdm(X_Test['Cleaned_text'].values):
2     list_of_sent_test_avg2v.append(sent_test_avg2v.split())
3
4 for sent_cv_avg2v in tqdm(X_CV['Cleaned_text'].values):
5     list_of_sent_cv_avg2v.append(sent_cv_avg2v.split())

```

100% |██████████| 20000/20000 [00:00<00:00, 152319.91it/s]
 100% |██████████| 16000/16000 [00:00<00:00, 189355.89it/s]

```

1 w2v_model_train = Word2Vec(list_of_sent_train_avg2v,min_count=5,size=50,workers=4)
2 w2v_words_svm_train=list(w2v_model_train.wv.vocab)

```

```

1 w2v_model_test = Word2Vec(list_of_sent_test_avg2v,min_count=5,size=50,workers=4)
2 w2v_words_svm_test=list(w2v_model_test.wv.vocab)

```

```

1 w2v_model_cv = Word2Vec(list_of_sent_cv_avg2v,min_count=5,size=50,workers=4)
2 w2v_words_svm_cv=list(w2v_model_cv.wv.vocab)

```

```

1 train_vectors=[];
2 for sent in list_of_sent_train_avg2v:
3     sent_vec=np.zeros(50)
4     cnt_words=0;
5     for word in sent:
6         if word in w2v_words_svm_train:
7             vec=w2v_model_train.wv[word]
8             sent_vec+=vec
9             cnt_words+=1
10    if cnt_words !=0:
11        sent_vec/=cnt_words
12    train_vectors.append(sent_vec)
13 print(len(train_vectors))
14 print(len(train_vectors[0]))

```

64000
50

```

1 test_vectors=[];
2 for sent in tqdm(list_of_sent_test_avg2v):
3     sent_vec=np.zeros(50)
4     cnt_words=0;
5     for word in sent:
6         if word in w2v_words_svm_test:
7             vec=w2v_model_test.wv[word]
8             sent_vec+=vec
9             cnt_words+=1
10    if cnt_words !=0:
11        sent_vec/=cnt_words
12    test_vectors.append(sent_vec)
13 print(len(test_vectors))
14 print(len(test_vectors[0]))

```

100% |██████████| 20000/20000 [00:16<00:00, 1200.95it/s]20000
50

```

1 cv_vectors=[];
2 for sent in tqdm(list_of_sent_cv_avgw2v):
3     sent_vec=np.zeros(50)
4     cnt_words=0;
5     for word in sent:
6         if word in w2v_words_svm_cv:
7             vec=w2v_model_cv.wv[word]
8             sent_vec+=vec
9             cnt_words+=1
10    if cnt_words !=0:
11        sent_vec/=cnt_words
12    cv_vectors.append(sent_vec)
13 print(len(cv_vectors))
14 print(len(cv_vectors[0]))

```


100%|██████████| 16000/16000 [00:12<00:00, 1257.27it/s]16000
50

```

1 results,best_alpha,linear_svm = SVM_Linear_SVM(train_vectors,Y_Train,'l2')
2 results.head()

```

0.8893952681617303
{'alpha': 0.001}



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_score	split1
0	0.464100	0.050165	0.035754	0.003152	0.0001	{'alpha': 0.0001}	0.887549	
1	0.268217	0.012332	0.034349	0.000384	0.001	{'alpha': 0.001}	0.889850	
2	0.218326	0.003064	0.034603	0.000560	0.01	{'alpha': 0.01}	0.890220	
3	0.218255	0.016672	0.035085	0.000583	0.1	{'alpha': 0.1}	0.888572	
4	0.207437	0.010266	0.035847	0.000986	1	{'alpha': 1}	0.880606	

▼ BEST ALPHA = 0.001 WITH AUC = 0.8890 ON TRAINING DATA

```
1 best_alpha = Find_best_Alpha(best_alpha)
```

0.001

```


1 linear_svm = SGDClassifier(loss='hinge',penalty='l2',random_state=None, class_weight=None,alpha=best_alpha)
2 clf=linear_svm.fit(train_vectors,Y_Train)

```

```

3 calibrated_model=CalibratedClassifierCV(clf,cv='prefit')
4 model_m=calibrated_model.fit(train_vectors,Y_Train)
5 pred_test_data=linear_svm.predict(test_vectors)
6 y_train_predicted_prob = model_m.predict_proba(train_vectors)[:,-1]
7 y_test_predicted_prob=model_m.predict_proba(test_vectors)[:,-1]
8 train_fpr, train_tpr, train_thresholds=roc_curve(Y_Train,y_train_predicted_prob,pos_label='positive')
9 test_fpr, test_tpr, test_thresholds = roc_curve(Y_Test, y_test_predicted_prob,pos_label='positive')
10 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
11 plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
12 plt.legend()
13 plt.xlabel("Alpha: hyperparameter")
14 plt.ylabel("AUC")
15 plt.title("ERROR PLOTS")
16 plt.grid()
17 plt.show()

```



TEST AUC = 0.80326 FOR ALPHA = 0.001

▼ SVM on Amazon Fine Food RReview using L1 Regularization

```


1 results,best_alpha,linear_svm = SVM_Linear_SVM(train_vectors,Y_Train,'l1')
2 results.head()

```



0.8868313361702558

{ 'alpha': 0.001}



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_score	split1
0	0.615362	0.034282	0.029011	0.000364	0.0001	{'alpha': 0.0001}	0.888781	
1	0.338329	0.021439	0.029690	0.001517	0.001	{'alpha': 0.001}	0.887724	
2	0.306859	0.007375	0.029253	0.000761	0.01	{'alpha': 0.01}	0.830911	
3	0.272682	0.004794	0.028153	0.000324	0.1	{'alpha': 0.1}	0.500000	
4	0.279753	0.009730	0.029353	0.001307	1	{'alpha': 1}	0.500000	

TF-IDF Word2Vec Vectorization Technique for SUPPORT VECTOR MACHINE on Amazon Review

```
1 model_Avgw2v = TfidfVectorizer()
2 X_Train_Avgw2v=model_Avgw2v.fit_transform(X_Train['Cleaned_text'].values)
```

```
1 X_Test_Avgw2v=model_Avgw2v.transform(X_Test['Cleaned_text'].values)
2 X_CV_Avgw2v=model_Avgw2v.transform(X_CV['Cleaned_text'].values)
```

```
1 dictionary = dict(zip(model_Avgw2v.get_feature_names(), list(model_Avgw2v.idf_)))
```

```
1 tfidf_feature=model_Avgw2v.get_feature_names()
2
3 tfidf_sent_vectors_train=[];
4 #final_tf_idf = [];
5 row=0;
6
7 for sent in tqdm(list_of_sent_train_avgw2v):
8     sent_vec=np.zeros(50)
9     weight_sum=0;
10    for word in sent :
11        if word in w2v_words_svm_train and word in tfidf_feature :
12            vec=w2v_model_train.wv[word]
13            #tf_idf=final_tf_idf[row,tfidf_feature.index(word)]
14            tf_idf=dictionary[word]*(sent.count(word)/len(sent))
15            sent_vec+=(vec*tf_idf)
```

```

16         weight_sum+=tf_idf
17
18     if weight_sum!=0:
19         sent_vec/=weight_sum
20     tfidf_sent_vectors_train.append(sent_vec)
21     row+=1

```

100% |██████████| 64000/64000 [14:28<00:00, 73.65it/s]

```

1 tfidf_sent_vectors_test=[];
2 #final_tf_idf = [];
3 row=0;
4
5 for sent in tqdm(list_of_sent_test_avgw2v):
6     sent_vec=np.zeros(50)
7     weight_sum=0;
8     for word in sent :
9         if word in w2v_words_svm_test and word in tfidf_feature :
10             vec=w2v_model_test.wv[word]
11             #tf_idf=final_tf_idf[row,tfidf_feature.index(word)]
12             tf_idf=dictionary[word]*(sent.count(word)/len(sent))
13             sent_vec+=(vec*tf_idf)
14             weight_sum+=tf_idf
15
16     if weight_sum!=0:
17         sent_vec/=weight_sum
18     tfidf_sent_vectors_test.append(sent_vec)
19     row+=1

```

100% |██████████| 20000/20000 [04:23<00:00, 75.82it/s]

```


1 results,best_alpha,linear_svm = SVM_Linear_SVM(tfidf_sent_vectors_train,Y_Train,'l2')
2 results.head()

```

↗

0.8593906637211086

{ 'alpha': 0.001}



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_score	split1
0	0.593249	0.028793	0.035492	0.001220	0.0001	{ 'alpha': 0.0001}	0.848499	
1	0.310280	0.012341	0.035850	0.002070	0.001	{ 'alpha': 0.001}	0.856023	
2	0.234241	0.006350	0.035853	0.001414	0.01	{ 'alpha': 0.01}	0.855193	
3	0.230646	0.004971	0.034734	0.001037	0.1	{ 'alpha': 0.1}	0.842829	
4	0.212810	0.006187	0.035471	0.001092	1	{ 'alpha': 1}	0.854022	

▼ BEST ALPHA = 0.001 WITH AUC = 0.8593 ON TRAINING DATA

```
1 best_alpha = Find_best_Alpha(best_alpha)
```


```
↳ 0.001
```

```

1 linear_svm = SGDClassifier(loss='hinge',penalty='l2',random_state=None, class_weight=None,alpha=best_alpha)
2 clf=linear_svm.fit(tfidf_sent_vectors_train,Y_Train)
3 calibrated_model=CalibratedClassifierCV(clf,cv='prefit')
4 model_m=calibrated_model.fit(tfidf_sent_vectors_train,Y_Train)
5 pred_test_data=linear_svm.predict(tfidf_sent_vectors_test)
6 y_train_predicted_prob = model_m.predict_proba(tfidf_sent_vectors_train)[: ,1]
7 y_test_predicted_prob=model_m.predict_proba(tfidf_sent_vectors_test)[: ,1]
8 train_fpr, train_tpr, train_thresholds=roc_curve(Y_Train,y_train_predicted_prob,pos_label='positive')
9 test_fpr, test_tpr, test_thresholds = roc_curve(Y_Test, y_test_predicted_prob,pos_label='positive')
10 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
11 plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
12 plt.legend()
13 plt.xlabel("Alpha: hyperparameter")
14 plt.ylabel("AUC")
15 plt.title("ERROR PLOTS")
16 plt.grid()
17 plt.show()


```

```
↳
```



```
1 results,best_alpha,linear_svm = SVM_Linear_SVM(tfidf_sent_vectors_train,Y_Train,'l1')
2 results.head()
```

```
0.8519893373085686
{'alpha': 0.0001}
```



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_score	split1
0	0.788327	0.064601	0.029334	0.000754	0.0001	{'alpha': 0.0001}	0.848334	
1	0.390466	0.020528	0.029376	0.001168	0.001	{'alpha': 0.001}	0.843245	
2	0.308332	0.004531	0.029242	0.000428	0.01	{'alpha': 0.01}	0.618617	
3	0.272443	0.007234	0.027860	0.000470	0.1	{'alpha': 0.1}	0.500000	
4	0.271363	0.005789	0.028100	0.000459	1	{'alpha': 1}	0.500000	

▼ IMPLEMENT SUPPORT VECTOR MACHINE USING RBF KERNEL

▼ SPLIT THE DATA INTO TRAIN , TEST & CV

```
1 from sklearn.model_selection import train_test_split
2 from sklearn.neighbors import KNeighborsClassifier
3 from sklearn.metrics import accuracy_score
4 from sklearn.model_selection import cross_val_score
```

```

5 from collections import Counter
6 from sklearn.metrics import confusion_matrix
7 from sklearn.metrics import classification_report
8
9 X_1,X_Test_RBF,Y_1,Y_Test_RBF = train_test_split(final_data_40K,amazon_polarity_labels_40K,test_size=0.2,random_st
10 X_Train_RBF,X_CV_RBF,Y_Train_RBF,Y_CV_RBF = train_test_split(X_1,Y_1,test_size=0.2)

```

```

1 print(X_Train_RBF.shape, Y_Train_RBF.shape)
2 print(X_CV_RBF.shape, Y_CV_RBF.shape)
3 print(X_Test_RBF.shape, Y_Test_RBF.shape)
4
5 print(""*100)
6
7
8 count_vector=CountVectorizer(min_df=10,max_features=500)
9 X_Train_data_bow_rbf=(count_vector.fit_transform(X_Train_RBF['Cleaned_text'].values))
10 X_Test_data_bow_rbf=(count_vector.transform(X_Test_RBF['Cleaned_text'].values))
11 X_CV_data_bow_rbf=(count_vector.transform(X_CV_RBF['Cleaned_text'].values))
12
13 print("After vectorizations")
14 print(X_Train_data_bow_rbf.shape, Y_Train_RBF.shape)
15 print(X_CV_data_bow_rbf.shape, Y_CV_RBF.shape)
16 print(X_Test_data_bow_rbf.shape, Y_Test_RBF.shape)
17 print(""*100)

```

```

↳ (19200, 12) (19200,)
   (4800, 12) (4800,)
   (6000, 12) (6000,)
=====
After vectorizations
(19200, 500) (19200,)
(4800, 500) (4800,)
(6000, 500) (6000,)
=====

```

```


1 from sklearn.model_selection import GridSearchCV
2 from scipy.stats import randint as sp_randint
3 from sklearn.model_selection import cross_val_score
4 from sklearn.svm import SVC
5
6 def RBF_SVM(x_training_data,y_training_data):
7     grid_params = { 'C' : [10**x for x in range(-4,5)]
8                     }
9     rbf_SVM = SVC(kernel='rbf',random_state=None, class_weight=None)
10    clf=GridSearchCV(rbf_SVM,grid_params,scoring='roc_auc',return_train_score=True,cv=3)
11    clf.fit(x_training_data,y_training_data)
12    results = pd.DataFrame.from_dict(clf.cv_results_)
13    results = results.sort_values(['param_C'])
14    train_auc= results['mean_train_score']
15    train_auc_std= results['std_train_score']
16    cv_auc = results['mean_test_score']
17    cv_auc_std= results['std_test_score']
18    best_c = results['param_C']
19    print(clf.best_score_)
20    print(clf.best_params_)
21    plt.plot(best_c, train_auc, label='Train AUC')
22    plt.plot(best_c, cv_auc, label='CV AUC')
23    plt.scatter(best_c, train_auc, label='Train AUC points')
24    plt.scatter(best_c, cv_auc, label='CV AUC points')
25    plt.legend()
26    plt.xlabel("C = 1/Alpha : hyperparameter")
27    plt.ylabel("AUC")
28    plt.title("Hyper parameter Vs AUC plot")
29    plt.grid()

```

```
30 plt.show()
31 return results,clf,rbf_SVM
```

```
1 results,best_one_by_alpha,rbf_svm = RBF_SVM(X_Train_data_bow_rbf,Y_Train_RBF)
2 results.head()
```

```
0.8744640078013962
{'C': 1}
```



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params	split0_test_score	split1_test_score
0	8.881859	0.075935	3.833624	0.016110	0.0001	{'C': 0.0001}	0.726542	0
1	11.182248	0.074443	4.521126	0.398858	0.001	{'C': 0.001}	0.841201	0
2	16.277469	0.170417	6.572980	0.049068	0.01	{'C': 0.01}	0.872069	0
3	16.762876	0.292377	7.022699	0.063776	0.1	{'C': 0.1}	0.872636	0
4	16.579718	0.298460	7.122815	0.083631	1	{'C': 1}	0.873039	0


```
1 def Find_best_C(best_c) :
2     best_C = best_c.best_params_
3     best_C=best_C.get("C")
4     print(best_C)
5     return best_C
```

```
1 best_c = Find_best_C(best_one_by_alpha)
```

```
1
```

```
1 rbf_svm = SVC(C=best_c, kernel='rbf',random_state=None, class_weight=None,probability=True)
2 rbf_svm.fit(X_Train_data_bow_rbf,Y_Train_RBF)
3 #calibrated_model=CalibratedClassifierCV(clf,cv='prefit')
4 #model_m=calibrated_model.fit(tfidf_sent_vectors_train,Y_Train)
5 pred_test_data=rbf_svm.predict(X_Test_data_bow_rbf)
6 y_train_predicted_prob = rbf_svm.predict_proba(X_Train_data_bow_rbf)[: ,1]
7 y_test_predicted_prob=rbf_svm.predict_proba(X_Test_data_bow_rbf)[: ,1]
8 train_fpr, train_tpr, train_thresholds=roc_curve(Y_Train_RBF,y_train_predicted_prob,pos_label='positive')
9 test_fpr, test_tpr, test_thresholds = roc_curve(Y_Test_RBF, y_test_predicted_prob,pos_label='positive')
10 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
11 plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
12 plt.legend()
13 plt.xlabel("C=1/Alpha: hyperparameter")
14 plt.ylabel("AUC")
15 plt.title("ERROR PLOTS")
```

```
16 plt.grid()
17 plt.show()
```



```
1 from sklearn.metrics import roc_auc_score
2 from sklearn.metrics import classification_report, confusion_matrix
3
4 roc_auc_score(Y_Test_RBF, y_test_predicted_prob)
```

```
0.892050278061514
```

```
1 print(classification_report(Y_Test_RBF, pred_test_data))
2 print(confusion_matrix(Y_Test_RBF, pred_test_data))
```

```
precision    recall  f1-score   support

negative     0.87     0.14     0.24         660
positive     0.90     1.00     0.95        5340

accuracy                   0.90        6000
macro avg     0.89     0.57     0.60        6000
weighted avg  0.90     0.90     0.87        6000

[[ 93 567]
 [ 14 5326]]
```

▼ SUPPORT VECTOR MACHINE TFIDF VECTORIZATION TECHNIQUE FOR RBF KI

```
1 from sklearn.feature_extraction.text import TfidfVectorizer
2
3 print(X_Train_RBF.shape, Y_Train_RBF.shape)
4 print(X_CV_RBF.shape, Y_CV_RBF.shape)
5 print(X_Test_RBF.shape, Y_Test_RBF.shape)
6
7 print("="*100)
8
9
10 tfidf_vector=TfidfVectorizer(min_df=10,max_features=500)
11 X_Train_data_tfidf=(tfidf_vector.fit_transform(X_Train_RBF['Cleaned_text'].values))
12 X_Test_data_tfidf=(tfidf_vector.transform(X_Test_RBF['Cleaned_text'].values))
13 X_CV_data_tfidf=(tfidf_vector.transform(X_CV_RBF['Cleaned_text'].values))
14
15 print("After vectorizations")
16 print(X_Train_data_tfidf.shape, Y_Train_RBF.shape)
17 print(X_CV_data_tfidf.shape, Y_CV_RBF.shape)
18 print(X_Test_data_tfidf.shape, Y_Test_RBF.shape)
19 print("="*100)
```

```

↳ (19200, 12) (19200,)
(4800, 12) (4800,)
(6000, 12) (6000,)
=====

```

```

After vectorizations
(19200, 500) (19200,)
(4800, 500) (4800,)
(6000, 500) (6000,)
=====

```

```


1 results,best_one_by_alpha,rbf_svm = RBF_SVM(X_Train_data_tfidf,Y_Train_RBF)
2 results.head()

```

```

↳ 0.8933775480843108
{'C': 10}

```



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params	split0_test_score	split1_test
0	8.389556	0.094299	3.762124	0.025662	0.0001	{'C': 0.0001}	0.848933	0
1	11.170771	0.033705	4.209851	0.006669	0.001	{'C': 0.001}	0.882891	0
2	15.728408	0.163814	5.994792	0.052254	0.01	{'C': 0.01}	0.888252	0
3	15.862320	0.210091	6.360304	0.062101	0.1	{'C': 0.1}	0.888043	0
4	15.695362	0.228676	6.734292	0.308750	1	{'C': 1}	0.889058	0

```

1 best_c = Find_best_C(best_one_by_alpha)

```

```

↳ 10

```

```

1 rbf_svm = SVC(C=best_c, kernel='rbf',random_state=None, class_weight=None,probability=True)
2 rbf_svm.fit(X_Train_data_tfidf,Y_Train_RBF)
3 #calibrated_model=CalibratedClassifierCV(clf,cv='prefit')
4 #model_m=calibrated_model.fit(tfidf_sent_vectors_train,Y_Train)
5 pred_test_data=rbf_svm.predict(X_Test_data_tfidf)
6 y_train_predicted_prob = rbf_svm.predict_proba(X_Train_data_tfidf)[:,-1]
7 y_test_predicted_prob=rbf_svm.predict_proba(X_Test_data_tfidf)[:,-1]
8 train_fpr, train_tpr, train_thresholds=roc_curve(Y_Train_RBF,y_train_predicted_prob,pos_label='positive')
9 test_fpr, test_tpr, test_thresholds = roc_curve(Y_Test_RBF, y_test_predicted_prob,pos_label='positive')
10 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
11 plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
12 plt.legend()
13 plt.xlabel("C=1/Alpha: hyperparameter")
14 plt.ylabel("AUC")

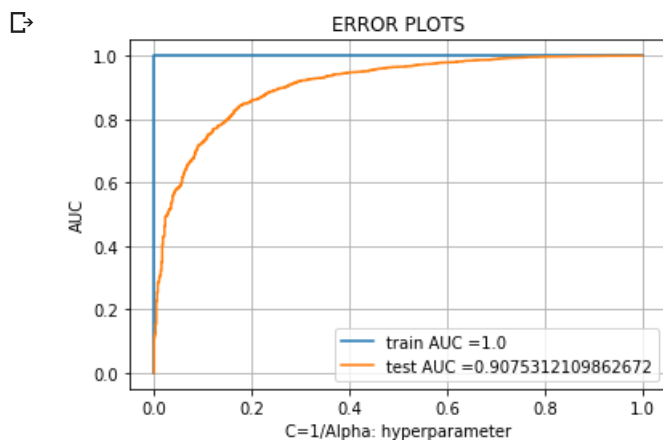
```



```

15 plt.title("ERROR PLOTS")
16 plt.grid()
17 plt.show()

```



```
1 roc_auc_score(Y_Test_RBF,y_test_predicted_prob)
```

```
0.9075312109862672
```

```

1 print(classification_report(Y_Test_RBF,pred_test_data))
2 print(confusion_matrix(Y_Test_RBF,pred_test_data))

```

```

precision    recall  f1-score   support

negative     0.73     0.37     0.49         660
positive     0.93     0.98     0.95        5340

accuracy          0.92         6000
macro avg         0.83     0.67     0.72         6000
weighted avg      0.90     0.92     0.90         6000

[[ 241  419]
 [  88 5252]]

```

▼ AVERAGE WORD2VEC VECTORIZATION TECHNIQUE USING RFB KERNEL

```

1 from gensim.models import Word2Vec
2 from gensim.models import KeyedVectors
3 import pickle
4
5 list_of_sent_train_avgw2v=[]
6 list_of_sent_test_avgw2v=[]
7 list_of_sent_cv_avgw2v=[]
8 for sent_train_avgw2v in tqdm(X_Train_RBF['Cleaned_text'].values):
9     list_of_sent_train_avgw2v.append(sent_train_avgw2v.split())

```

```
100%|██████████| 19200/19200 [00:00<00:00, 170831.10it/s]
```

```

1 for sent_test_avgw2v in tqdm(X_Test_RBF['Cleaned_text'].values):
2     list_of_sent_test_avgw2v.append(sent_test_avgw2v.split())
3
4 for sent_cv_avgw2v in tqdm(X_CV_RBF['Cleaned_text'].values):
5     list_of_sent_cv_avgw2v.append(sent_cv_avgw2v.split())

```

```

100%|██████████| 6000/6000 [00:00<00:00, 198876.43it/s]
100%|██████████| 4800/4800 [00:00<00:00, 165923.50it/s]

```

```
1 w2v_model_train = Word2Vec(list_of_sent_train_avgw2v,min_count=5,size=50,workers=4)
2 w2v_words_svmbf_train=list(w2v_model_train.wv.vocab)
```

```
1 w2v_model_test = Word2Vec(list_of_sent_test_avgw2v,min_count=5,size=50,workers=4)
2 w2v_words_svmbf_test=list(w2v_model_test.wv.vocab)
```

```
1 w2v_model_cv = Word2Vec(list_of_sent_cv_avgw2v,min_count=5,size=50,workers=4)
2 w2v_words_svmbf_cv=list(w2v_model_cv.wv.vocab)
```

```
1 train_vectors=[];
2 for sent in list_of_sent_train_avgw2v:
3     sent_vec=np.zeros(50)
4     cnt_words=0;
5     for word in sent:
6         if word in w2v_words_svmbf_train:
7             vec=w2v_model_train.wv[word]
8             sent_vec+=vec
9             cnt_words+=1
10    if cnt_words !=0:
11        sent_vec/=cnt_words
12    train_vectors.append(sent_vec)
13 print(len(train_vectors))
14 print(len(train_vectors[0]))
```

```
19200
50
```


```
1 test_vectors=[];
2 for sent in list_of_sent_test_avgw2v:
3     sent_vec=np.zeros(50)
4     cnt_words=0;
5     for word in sent:
6         if word in w2v_words_svmbf_test:
7             vec=w2v_model_test.wv[word]
8             sent_vec+=vec
9             cnt_words+=1
10    if cnt_words !=0:
11        sent_vec/=cnt_words
12    test_vectors.append(sent_vec)
13 print(len(test_vectors))
14 print(len(test_vectors[0]))
```

```
6000
50
```

```
1 cv_vectors=[];
2 for sent in list_of_sent_cv_avgw2v:
3     sent_vec=np.zeros(50)
4     cnt_words=0;
5     for word in sent:
6         if word in w2v_words_svmbf_cv:
7             vec=w2v_model_cv.wv[word]
8             sent_vec+=vec
9             cnt_words+=1
10    if cnt_words !=0:
11        sent_vec/=cnt_words
12    cv_vectors.append(sent_vec)
13 print(len(cv_vectors))
14 print(len(cv_vectors[0]))
```

```
1 results,best_one_by_alpha,rbf_svm = RBF_SVM(train_vectors,Y_Train_RBF)
2 results.head()
```

```
0.854387844023114
{'C': 10}
```




	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params	split0_test_score	split1_test_score
0	4.219129	0.043246	1.662449	0.004553	0.0001	{'C': 0.0001}	0.798200	0
1	4.590138	0.026659	1.666605	0.004292	0.001	{'C': 0.001}	0.844461	0
2	5.761482	0.085106	1.824312	0.001734	0.01	{'C': 0.01}	0.851311	0
3	5.973886	0.087304	2.019669	0.227633	0.1	{'C': 0.1}	0.853838	0
4	5.691197	0.070469	1.847777	0.007446	1	{'C': 1}	0.854072	0

```
1 best_c = Find_best_C(best_one_by_alpha)
```

```
10
```

```
1 rbf_svm = SVC(C=best_c, kernel='rbf',random_state=None, class_weight=None,probability=True)
2 rbf_svm.fit(train_vectors,Y_Train_RBF)
3 #calibrated_model=CalibratedClassifierCV(clf,cv='prefit')
4 #model_m=calibrated_model.fit(tfidf_sent_vectors_train,Y_Train)
5 pred_test_data=rbf_svm.predict(test_vectors)
6 y_train_predicted_prob = rbf_svm.predict_proba(train_vectors)[: ,1]
7 y_test_predicted_prob=rbf_svm.predict_proba(test_vectors)[: ,1]
8 train_fpr, train_tpr, train_thresholds=roc_curve(Y_Train_RBF,y_train_predicted_prob,pos_label='positive')
9 test_fpr, test_tpr, test_thresholds = roc_curve(Y_Test_RBF, y_test_predicted_prob,pos_label='positive')
10 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
11 plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
12 plt.legend()
13 plt.xlabel("C=1/Alpha: hyperparameter")
14 plt.ylabel("AUC")
15 plt.title("ERROR PLOTS")
16 plt.grid()
17 plt.show()
```

```
10
```



```
1 from sklearn.metrics import roc_auc_score
2 from sklearn.metrics import classification_report, confusion_matrix
3
4 roc_auc_score(Y_Test_RBF, y_test_predicted_prob)
```

```
0.6155249120417661
```

```
1 print(classification_report(Y_Test_RBF, pred_test_data))
2 print(confusion_matrix(Y_Test_RBF, pred_test_data))
```

```
precision    recall  f1-score   support

negative     0.14     0.00     0.00         660
positive     0.89     1.00     0.94        5340

accuracy              0.89        6000
macro avg     0.52     0.50     0.47        6000
weighted avg     0.81     0.89     0.84        6000

[[ 1 659]
 [ 6 5334]]
```

▼ TF-IDF Average Word2Vec using RBF Kernel

```
1 model_Avgw2v = TfidfVectorizer()
2 X_Train_Avgw2v=model_Avgw2v.fit_transform(X_Train_RBF['Cleaned_text'].values)
```

```
1 X_Test_Avgw2v=model_Avgw2v.transform(X_Test_RBF['Cleaned_text'].values)
2 X_CV_Avgw2v=model_Avgw2v.transform(X_CV_RBF['Cleaned_text'].values)
```

```
1 dictionary = dict(zip(model_Avgw2v.get_feature_names(), list(model_Avgw2v.idf_)))
```

```
1 tfidf_feature=model_Avgw2v.get_feature_names()
2
3 tfidf_sent_vectors_train=[];
4 #final_tf_idf = [];
5 row=0;
6
7 for sent in tqdm(list_of_sent_train_avgw2v):
8     sent_vec=np.zeros(50)
9     weight_sum=0;
10    for word in sent :
11        if word in w2v_words_svmbf_train and word in tfidf_feature :
```

```

12         vec=w2v_model_train.wv[word]
13         #tf_idf=final_tf_idf[row,tfidf_feature.index(word)]
14         tf_idf=dictionary[word]*(sent.count(word)/len(sent))
15         sent_vec+=(vec*tf_idf)
16         weight_sum+=tf_idf
17
18     if weight_sum!=0:
19         sent_vec/=weight_sum
20     tfidf_sent_vectors_train.append(sent_vec)
21     row+=1

```

100% |██████████| 19200/19200 [02:15<00:00, 141.21it/s]

```

1 tfidf_sent_vectors_test=[];
2 #final_tf_idf = [];
3 row=0;
4
5 for sent in tqdm(list_of_sent_test_avgw2v):
6     sent_vec=np.zeros(50)
7     weight_sum=0;
8     for word in sent :
9         if word in w2v_words_svmbf_test and word in tfidf_feature :
10             vec=w2v_model_test.wv[word]
11             #tf_idf=final_tf_idf[row,tfidf_feature.index(word)]
12             tf_idf=dictionary[word]*(sent.count(word)/len(sent))
13             sent_vec+=(vec*tf_idf)
14             weight_sum+=tf_idf
15
16     if weight_sum!=0:
17         sent_vec/=weight_sum
18     tfidf_sent_vectors_test.append(sent_vec)
19     row+=1

```

100% |██████████| 6000/6000 [00:40<00:00, 148.81it/s]

```


1 results,best_one_by_alpha,rbf_svm = RBF_SVM(tfidf_sent_vectors_train,Y_Train_RBF)
2 results.head()

```

↗

0.8291443099232637

{ 'C': 10 }



	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params	split0_test_score	split1_test
0	4.131723	0.030522	1.669118	0.010235	0.0001	{'C': 0.0001}	0.724419	0
1	4.467314	0.044934	1.671763	0.007991	0.001	{'C': 0.001}	0.784976	0
2	5.640827	0.061905	1.883818	0.011674	0.01	{'C': 0.01}	0.820578	0
3	6.628933	0.044802	1.990721	0.026210	0.1	{'C': 0.1}	0.830384	0
4	6.555834	0.098492	1.970353	0.011784	1	{'C': 1}	0.831317	0

```
1 best_c = Find_best_C(best_one_by_alpha)
```


```
↳ 10
```

```

1 rbf_svm = SVC(C=best_c, kernel='rbf', random_state=None, class_weight=None, probability=True)
2 rbf_svm.fit(tfidf_sent_vectors_train, Y_Train_RBF)
3 #calibrated_model=CalibratedClassifierCV(clf, cv='prefit')
4 #model_m=calibrated_model.fit(tfidf_sent_vectors_train, Y_Train)
5 pred_test_data=rbf_svm.predict(tfidf_sent_vectors_test)
6 y_train_predicted_prob = rbf_svm.predict_proba(tfidf_sent_vectors_train)[: ,1]
7 y_test_predicted_prob=rbf_svm.predict_proba(tfidf_sent_vectors_test)[: ,1]
8 train_fpr, train_tpr, train_thresholds=roc_curve(Y_Train_RBF, y_train_predicted_prob, pos_label='positive')
9 test_fpr, test_tpr, test_thresholds = roc_curve(Y_Test_RBF, y_test_predicted_prob, pos_label='positive')
10 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
11 plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
12 plt.legend()
13 plt.xlabel("C=1/Alpha: hyperparameter")
14 plt.ylabel("AUC")
15 plt.title("ERROR PLOTS")
16 plt.grid()
17 plt.show()

```

```
↳
```



```
1 roc_auc_score(Y_Test_RBF,y_test_predicted_prob)
```

```
0.5950913630688912
```

```
1 print(classification_report(Y_Test_RBF,pred_test_data))
2 print(confusion_matrix(Y_Test_RBF,pred_test_data))
```

```
precision    recall  f1-score   support

negative     0.16     0.01     0.01         660
positive     0.89     0.99     0.94        5340

accuracy                   0.89        6000
macro avg                 0.52     0.50     0.48        6000
weighted avg              0.81     0.89     0.84        6000

[[  5  655]
 [ 27 5313]]
```

▼ PRETTY TABLE

```
1 pip install -U PTable
```

```
Collecting PTable
  Downloading https://files.pythonhosted.org/packages/ab/b3/b54301811173ca94119eb474634f120a49cd370f257d1aae5a4a
Building wheels for collected packages: PTable
  Building wheel for PTable (setup.py) ... done
  Created wheel for PTable: filename=PTable-0.9.2-cp36-none-any.whl size=22908 sha256=34e773e9cf72ff9c2b6259e9be
  Stored in directory: /root/.cache/pip/wheels/22/cc/2e/55980bfe86393df3e9896146a01f6802978d09d7ebc5a5e56
Successfully built PTable
Installing collected packages: PTable
Successfully installed PTable-0.9.2
```

```
1 from prettytable import PrettyTable
2
3 x= PrettyTable()
4 x.field_names = ["Vectorizer" , "Hyperparameter(Alpha/C)", "AUC","Regularization","Kernel"]
5 x.add_row(["Bag Of Words",0.001,0.9266,'L2',"Linear SVM"])
6 x.add_row(["Bag Of Words",0.0001,0.8836,'L1',"Linear SVM"])
7 x.add_row(["Tf-Idf",0.0001,0.9353,"L2","Linear SVM"])
8 x.add_row(["Tf-Idf",0.0001,0.9180,"L1","Linear SVM"])
9 x.add_row(["Avg Word2Vec",0.001,0.8893,"L2","Linear SVM"])
10 x.add_row(["Avg Word2Vec",0.001,0.8868,"L1","Linear SVM"])
11 x.add_row(["TF-IDF Word2Vec",0.001,0.8593,"L2","Linear SVM"])
12 x.add_row(["TF-IDF Word2Vec",0.0001,0.8519,"L1","Linear SVM"])
```

```

13 x.add_row(["Bag Of Words ",1,0.8744,"N/A","RBF SVM"])
14 x.add_row(["TF-IDF ",10,0.8933,"N/A","RBF SVM"])
15 x.add_row(["Avg Word2Vec ",10,0.8543,"N/A","RBF SVM"])
16 x.add_row(["TF-IDF Word2Vec ",10,0.82914,"N/A","RBF SVM"])
17 print(x)

```

Vectorizer	Hyperparameter(Alpha/C)	AUC	Regularization	Kernel
Bag Of Words	0.001	0.9266	L2	Linear SVM
Bag Of Words	0.0001	0.8836	L1	Linear SVM
Tf-Idf	0.0001	0.9353	L2	Linear SVM
Tf-Idf	0.0001	0.918	L1	Linear SVM
Avg Word2Vec	0.001	0.8893	L2	Linear SVM
Avg Word2Vec	0.001	0.8868	L1	Linear SVM
TF-IDF Word2Vec	0.001	0.8593	L2	Linear SVM
TF-IDF Word2Vec	0.0001	0.8519	L1	Linear SVM
Bag Of Words	1	0.8744	N/A	RBF SVM
TF-IDF	10	0.8933	N/A	RBF SVM
Avg Word2Vec	10	0.8543	N/A	RBF SVM
TF-IDF Word2Vec	10	0.82914	N/A	RBF SVM

▼ CONCLUSION : IT HAS BEEN OBSERVED THAT

- 1) FOR RBF KERNEL , AS C INCREASES , THE AUC SCORE DECREASES ON TEST DATA
- 2) FOR LINEAR SVM , THERE IS SLIGHTLY DECREASE IN AUC SCORE
- 3)* FOR LINEAR SVM :THE BEST ALPHA = 0.0001 AND AUC SCORE = 0.9353 USING TF-IDF VECTORIZ TECHNIQUE USING L2 REGULARIZATION*
- 4) FOR RBF SVM , BEST C = 10 AND AUC SCORE = 0.8933 USING TF-IDF VECTORIZATION TECHNIQUE FOLD CROSS VALIDATION